C.1 Random selection ICL across all models

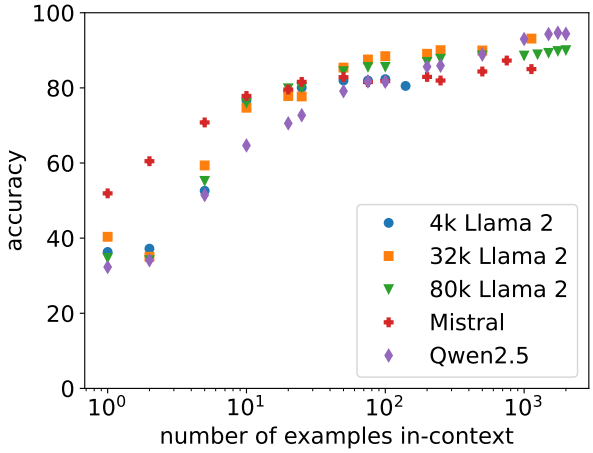

(a) TREC

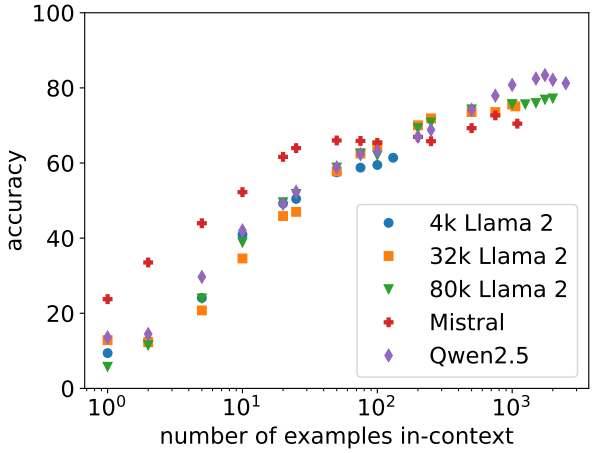

(b) TREC-fine

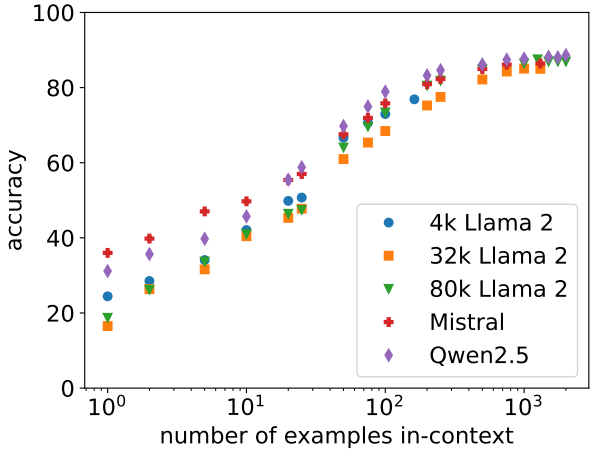

(c) NLU

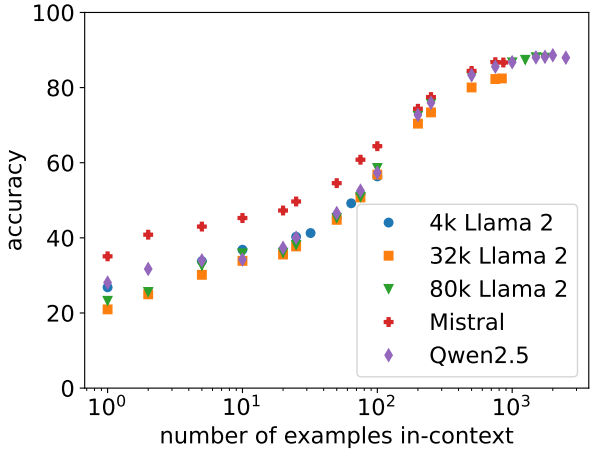

(d) Banking-77

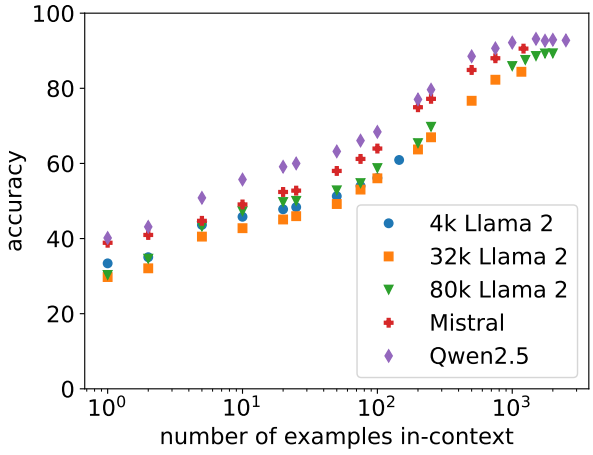

(e) Clinic-150

Figure 11: Performance of random-selection ICL across all models for each classification dataset. Performance continues to increase with additional examples in-context.

C.2 Retrieval ICL across all models

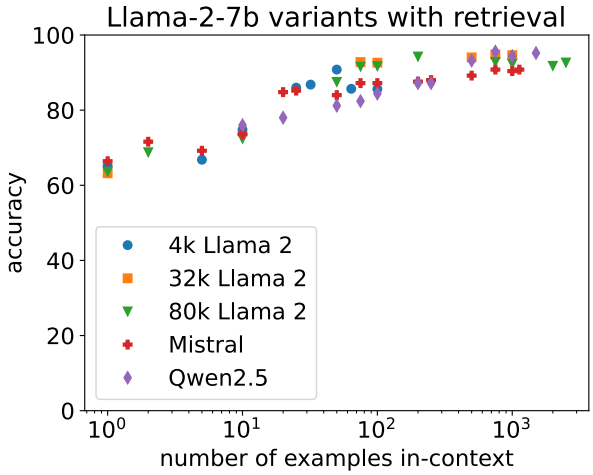

(a) TREC

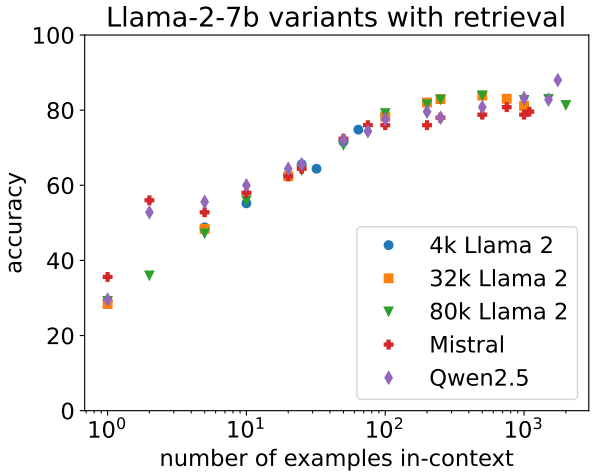

(b) TREC-fine

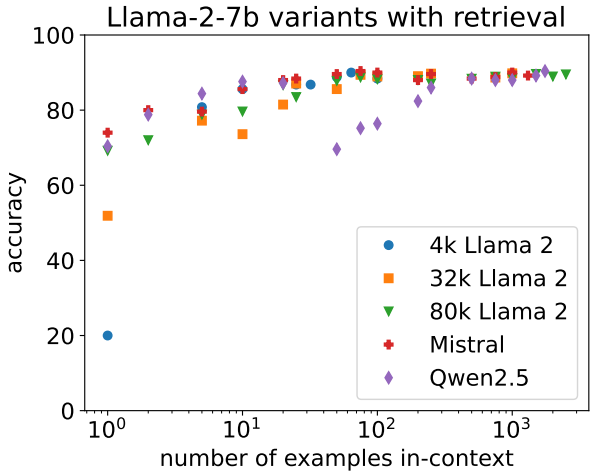

(c) NLU

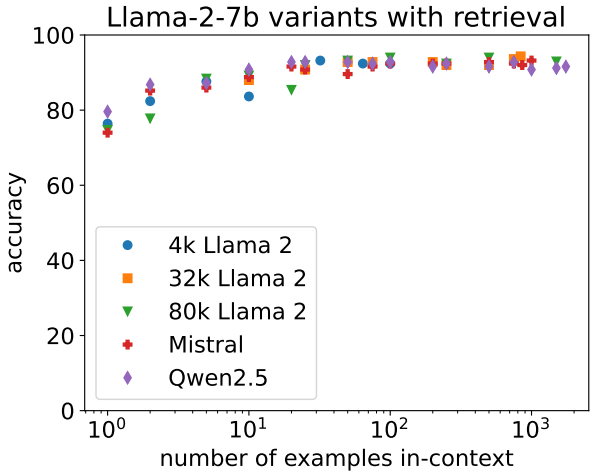

(d) Banking-77

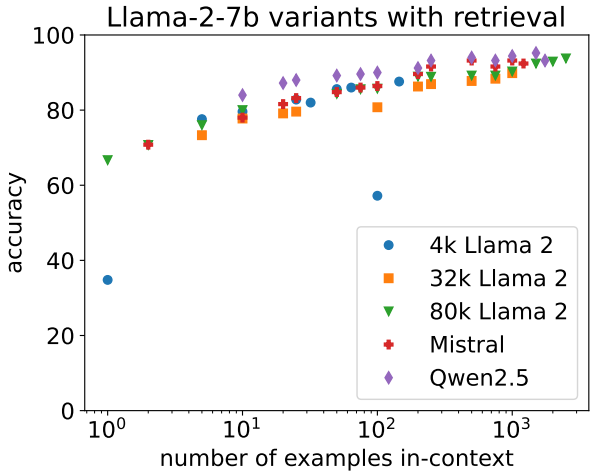

(e) Clinic-150

Figure 12: Performance of retrieval-based ICL across all models for each classification dataset. Short-context performance here is higher than for random-selection, but performance continues to improve with more examples until a saturation point, where performance flattens out.

C.3 Comparing retrieval, random selection, and finetuning

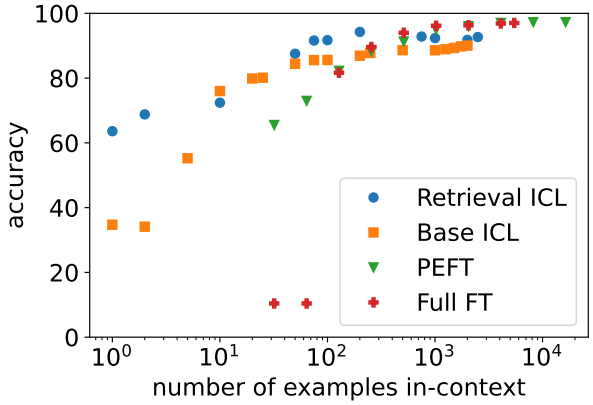

(a) TREC

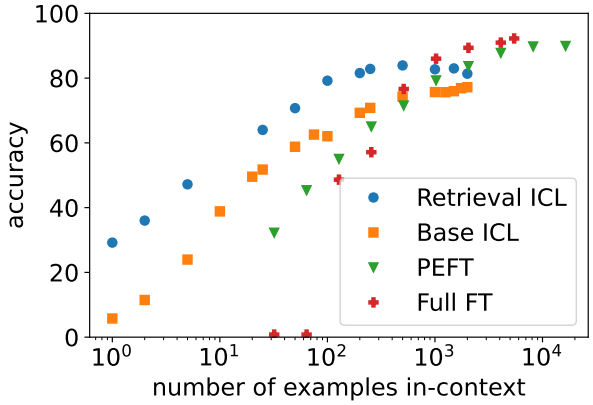

(b) TREC-fine

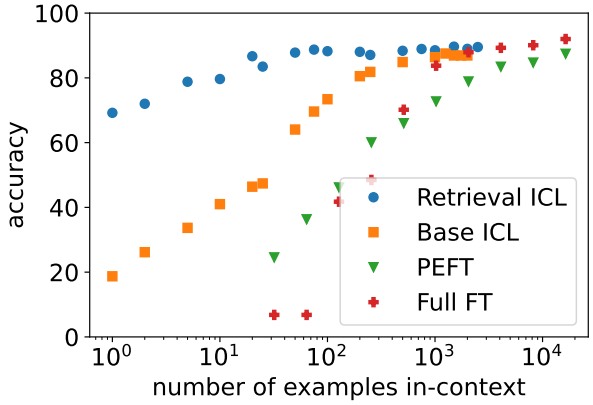

(c) NLU

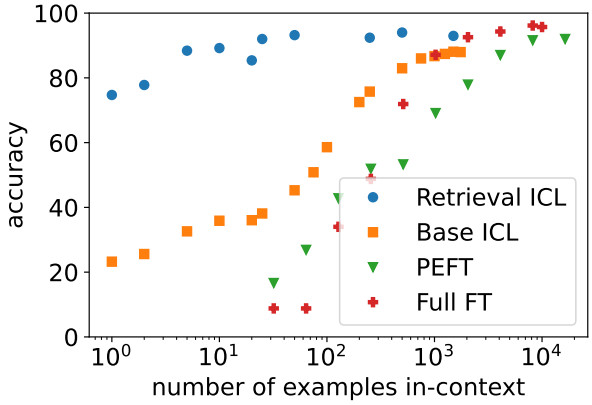

(d) Banking-77

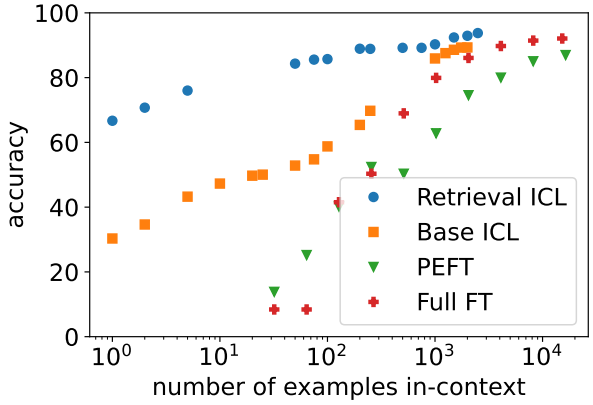

(e) Clinic-150

Figure 13: Performance of retrieval-based ICL, random-selection ICL, and finetuning across 5 datasets. At small example counts, ICL outperforms finetuning; when several thousand examples are used, finetuning outperforms ICL in some datasets.